

# EU AI Act: Draft Guidance on Reporting Serious AI Incidents

## 1. Executive Summary & Context

On September 26, 2025, the European Commission issued a draft Guidance to clarify the serious incident reporting obligations mandated by **Article 73 of the EU AI Act (Regulation (EU) 2024/1689)**.

- **Target Audience:** The obligation applies strictly to "**providers**" of high-risk AI systems (entities developing high-risk AI systems and placing them on the market under their own name/trademark).
- **Public Consultation:** The consultation period concluded on November 7, 2025.
- **Application Date:** The final version of this Guidance is expected to officially apply from **August 2, 2026**.

## 2. Core Concepts: Scope and Broad Definition

The Guidance establishes a remarkably broad interpretation of what triggers a reporting duty, particularly regarding causality.

### Defining a "Serious Incident" (Article 3(49))

A serious incident is defined as any malfunction or incident of an AI system that directly or indirectly results in one of the following serious outcomes:

- Death or serious harm to an individual's health.
- Serious and irreversible disruption to critical infrastructure management or operation.
- Infringements of obligations under EU law intended to protect fundamental rights.
- Serious harm to property or the environment.

### Key Legal Takeaway: Indirect Causality

The Commission explicitly notes that an **indirect causal link** between the AI system and the final harm is sufficient to trigger the duty to report.

Examples provided in the Guidance: \* **Healthcare:** An incorrect medical analysis by an AI system that leads to harm only after a subsequent clinical decision by a doctor still qualifies as a reportable incident. \* **Finance/HR:** A loan denial based on

a flawed AI assessment, or discriminatory treatment of applicants via an AI-enabled selection tool.

### 3. Strict Reporting Timelines

The draft guidance outlines highly compressed, tier-based notification windows based on the severity of the outcome.

| Severity / Scenario                                | Mandatory Reporting Deadline                                                                                                      |
|----------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| General Serious Incident                           | Without undue delay, and at most <b>within 15 days</b> of becoming aware.                                                         |
| Incidents Involving Death                          | Within <b>10 days</b> if a death may have been caused.                                                                            |
| Widespread Infringements / Critical Infrastructure | Within <b>2 days</b> for widespread fundamental rights violations or serious, irreversible disruption to critical infrastructure. |

Note: To accommodate these tight windows, Article 73(5) permits providers to initially submit an incomplete report and follow up with supplemental data later.

### 4. Post-Reporting & Investigative Duties

Filing a report triggers mandatory follow-up actions for the provider under Article 73(6):

- **Prompt Investigation:** Providers must immediately investigate the incident.
- **Evidence/System Preservation:** The Guidance explicitly prohibits providers from altering the AI system during an active investigation without first informing market surveillance authorities, provided those changes could impact the analysis of the incident.

## 5. Overlap with Other EU Legislation

The Guidance addresses how Article 73 interacts with pre-existing EU regulatory frameworks:

- **Sector-Specific Relief (e.g., NIS-2 Directive):** To prevent double-reporting in heavily regulated sectors (like critical infrastructure), a simplified regime is proposed. In these sectors, providers only report **fundamental rights violations** under the AI Act. Other technical/operational incidents should be handled via the sector-specific rules (e.g., NIS-2).
- **General Purpose AI (GPAI):** Article 55(1)(c) places a separate reporting duty on providers of GPAI models with systemic risks to notify the EU AI Office. While this specific Guidance does *not* formally cover GPAI, the Commission advises entities to align their internal workflows with both regimes.

## 6. Enforcement Risks and Non-Compliance Penalties

Failure to properly adhere to the incident reporting protocols exposes organizations to severe regulatory and civil liabilities:

- **Administrative Fines:** Up to **€15 million or 3% of worldwide annual turnover**, whichever is higher.
- **Market Surveillance Actions:** Within 7 days of a report, authorities can mandate aggressive corrective actions under the EU Market Surveillance Regulation (2019/1020), including product recalls, market withdrawals, or complete bans.
- **Supervisory Measures:** Public disclosure of enforcement actions, distribution restrictions, and official warnings.
- **Civil Claims:** Increased exposure to private civil litigations from affected end-users.

## 7. Recommended Action Points for Organizations

To prepare for the August 2026 rollout, the paper recommends that legal and compliance teams immediately integrate these expectations into their current structures:

- **Update Incident Response Protocols:** Build internal monitoring mechanisms specifically calibrated to detect "indirect harm" and fundamental rights risks.
- **Evidence Preservation Procedures:** Establish clear protocols regarding when system updates/patches must be frozen during a suspected incident.
- **Cross-Framework Alignment:** Ensure the AI incident workflow seamlessly maps onto existing GDPR data breach and product-safety frameworks.