

# Summary: A Robust Governance for the AI Act

## 1. Introduction & Core Ambitions

The paper addresses the critical challenge of enforcing the EU's new Artificial Intelligence Act (AIA) across complex multilevel governance systems (supranational and national). It aims to achieve two primary goals:

1. **Explain the current governance framework** established under Chapter VII of the AIA.
2. **Propose a normative "robust governance" model**—defined as the capacity to maintain core values, functions, and structures through adaptation and innovation when facing disruptive technological shifts.

## 2. Institutional Design Options

The authors evaluate how administrative structures can be organized at both the EU and national levels, highlighting three structural choices for each:

### Supranational (EU) Options

- **Option 1 (Centralized):** Absorb AI tasks directly within European Commission departments (e.g., Directorates-General like DG-CNECT). This maximizes Commission oversight but risks over-centralization.
- **Option 2 (Decentralized):** Establish a fully independent external EU AI agency at arm's length from the Commission. This limits daily Commission management.
- **Option 3 (Hybrid):** Retain a dedicated Commission DG working alongside arms-length EU-level agencies. **(Authors' Preference for Robustness).**

### National (Member State) Options

- **Option 1:** Create an entirely new, dedicated national AI agency. Attracts tailored skills but lacks industry-specific expertise and takes too long to legally establish.
- **Option 2:** Assign enforcement to existing sectoral agencies (e.g., banking or data protection authorities). Leverages existing infrastructure but risks narrow sectoral tunnel-vision.
- **Option 3:** Establish a centralized "**competence center**" within an experienced existing authority. **(Authors' Preference).** It temporarily or permanently pools interdisciplinary teams (e.g., lawyers and computer scientists) for specific cases, blending centralized tech expertise with sectoral insights.

### 3. The Current AIA Governance Architecture

The paper maps out the primary genetic (appointing) and functional (advisory/operational) linkages across three levels:

#### Supranational Level

- **European Commission:** Holds primary authority over non-legislative acts, including issuing guidelines, managing delegated acts, and running comitology procedures.
- **AI Office:** Established in January 2024 under a centralized design (Option 1) within **DG-CNECT**. It lacks a distinct legal personality and primarily issues non-binding support and advisory frameworks to the Commission, focusing heavily on General-Purpose AI (GPAI) oversight.
- **European AI Board:** Composed of representatives from each Member State. It coordinates national authorities and issues recommendations on implementing acts. Crucially, it **lacks binding power** to resolve cross-border disputes or overrule national agencies.
- **Advisory Forum & Scientific Panel:** The Forum consists of mixed commercial/civil society stakeholders to provide market insights. The Panel consists of independent academic experts who issue qualified alerts on GPAI systemic risks.

#### National & Corporate Level

- **Notifying Authorities:** National entities that designate, monitor, and accredit third-party Conformity Assessment Bodies.
- **Notified Bodies:** Independent third-party entities that perform the actual conformity testing and issue certificates for high-risk AI systems.
- **Market Surveillance Authorities (MSAs):** National single points of contact that investigate non-compliance, manage real-world testing, access source codes, and issue penalties.
- **Providers:** Corporate entities developing high-risk systems or GPAI models.

## 4. Key Legislative Tasks & Friction Points

| Regulatory Focus                 | Commission & AI Office Responsibilities                                                                                   | Identified Vulnerabilities & Friction Points                                                                                                                                                                                                                          |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Operationalizing Risk</b>     | Define rules for <b>"significant modifications"</b> that trigger a mandatory re-evaluation of high-risk systems.          | Standard model fine-tuning shouldn't count as a significant modification unless safety layers are removed or biases introduced. The authors champion medical-style <b>"change management plans"</b> to pre-approve dynamic updates safely without slowing deployment. |
| <b>GPAI Classification</b>       | Classify high-impact GPAI models as possessing <b>"systemic risk"</b> (using a presumptive threshold of $10^{25}$ FLOPs). | The $10^{25}$ FLOP threshold is outdated rapidly due to high-performing "smaller" models. Tech providers with deep pockets will likely exploit classification appeal mechanisms to delay stringent safety regulations.                                                |
| <b>Conformity &amp; Auditing</b> | Ensure impartiality, compliance, and prevent <b>"conformity shopping"</b> across jurisdictions.                           | Notified bodies rely on financial relationships (fees/commissions) with the very AI providers they audit, creating clear conflicts of interest. There is an extreme lack of operational transparency due to extensive outsourcing.                                    |
| <b>Enforcement Exemptions</b>    | Manage compliance timelines and transition periods.                                                                       | High-risk systems already on the market 24 months after entry into force are completely exempt from the AIA unless structurally modified. This "blank exemption" violates standard                                                                                    |

|  |  |                                                                         |
|--|--|-------------------------------------------------------------------------|
|  |  | product safety law principles and anti-competitively favors incumbents. |
|--|--|-------------------------------------------------------------------------|

## 5. The Authors' 5 Normative Recommendations

To streamline and patch the institutional deficits of the AIA, the paper outlines five core changes:

### A. Decentralize the AI Office

The AI Office should be detached from DG-CNECT and turned into an autonomous, decentralized agency with its own legal identity (modeled after the European Medicines Agency—EMA). This grants it the legal independence to issue binding decisions away from the political agendas of the Commission.

### B. Consolidate Advisory Bodies (Ockham's Razor)

The **Advisory Forum** and **Scientific Panel** should be merged into a single integrated advisory body. This eliminates bureaucratic redundancy, reduces administrative overhead, and blends commercial stakeholder feedback with independent scientific expertise to offer a single, unified voice.

### C. Create an AI Coordination Hub (EU AICH)

To prevent uncoordinated, fragmented enforcement across overlapping regulatory structures (e.g., the AI Office conflicting with the European Data Protection Board or national data authorities, like Italy's unilateral ban on ChatGPT), the EU should establish a centralized **AI Coordination Hub** to handle collective cross-agency decision-making.

### D. Empower the AI Board against Misuse

The AI Board must be granted the power to overrule or correct idiosyncratic national enforcement actions, mimicking the dispute resolution powers of the EDPB under the GDPR. Without this check, member states undergoing democratic backsliding could exploit ambiguities to establish illiberal biometric or facial surveillance regimes unchecked.

### E. Embed Institutional Learning Mechanisms

A dedicated unit should be established within the AI Office explicitly tasked with compiling and sharing "best and worst practices" across Member State competence centers. This builds active ex-ante and ex-post review cycles directly into the administrative culture to keep pace with rapid AI adjustments.